

Install the Necessary Libraries

Name : Pratik Yuvraj Yawalkar

Roll No. : AI23MTECH11006

Department : AI & ML

```
!pip install astroML numpy pandas scipy matplotlib seaborn

Collecting astroML
  Downloading astroML-1.0.2.post1-py3-none-any.whl (134 kB)
    134.3/134.3 kB 1.5 MB/s eta
0:00:00
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.23.5)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (1.5.3)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (1.11.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.10/dist-packages (from astroML) (1.2.2)
Requirement already satisfied: astropy>=3.0 in /usr/local/lib/python3.10/dist-packages (from astroML) (5.3.4)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.3.post1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.47.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (23.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.1)
Requirement already satisfied: pyerfa>=2.0 in /usr/local/lib/python3.10/dist-packages (from astropy>=3.0->astroML)
```

```
(2.0.1.1)
Requirement already satisfied: PyYAML>=3.13 in
/usr/local/lib/python3.10/dist-packages (from astropy>=3.0->astroML)
(6.0.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.1-
>pandas) (1.16.0)
Requirement already satisfied: joblib>=1.1.1 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18-
>astroML) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18-
>astroML) (3.2.0)
Installing collected packages: astroML
Successfully installed astroML-1.0.2.post1
```

1. In the class, we demonstrated the Central Limit Theorem for a sample drawn from a uniform distribution. Reproduce a similar plot for a sample drawn the from chi-square distribution with degrees of freedom equal to 3, for samples drawn once, 5 times, and 10 times. Either plot all of these on one multipanel figure similar to AstroML figure 3.20. (Hint: look up `numpy.random.chisquare` and show the distribution of `x` from 0 to 10)

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

def first_plot(sample_size: int =1)->None:

    np.random.seed(42) # Set the seed to draw the random samples
    df = 3             # Define the degrees of freedom for the chi-
square distribution
```

```

# Generate and plot samples for given Sample Size
# Draw samples from the chi-square distribution
samples = np.random.chisquare(df, size=(sample_size, 10000))

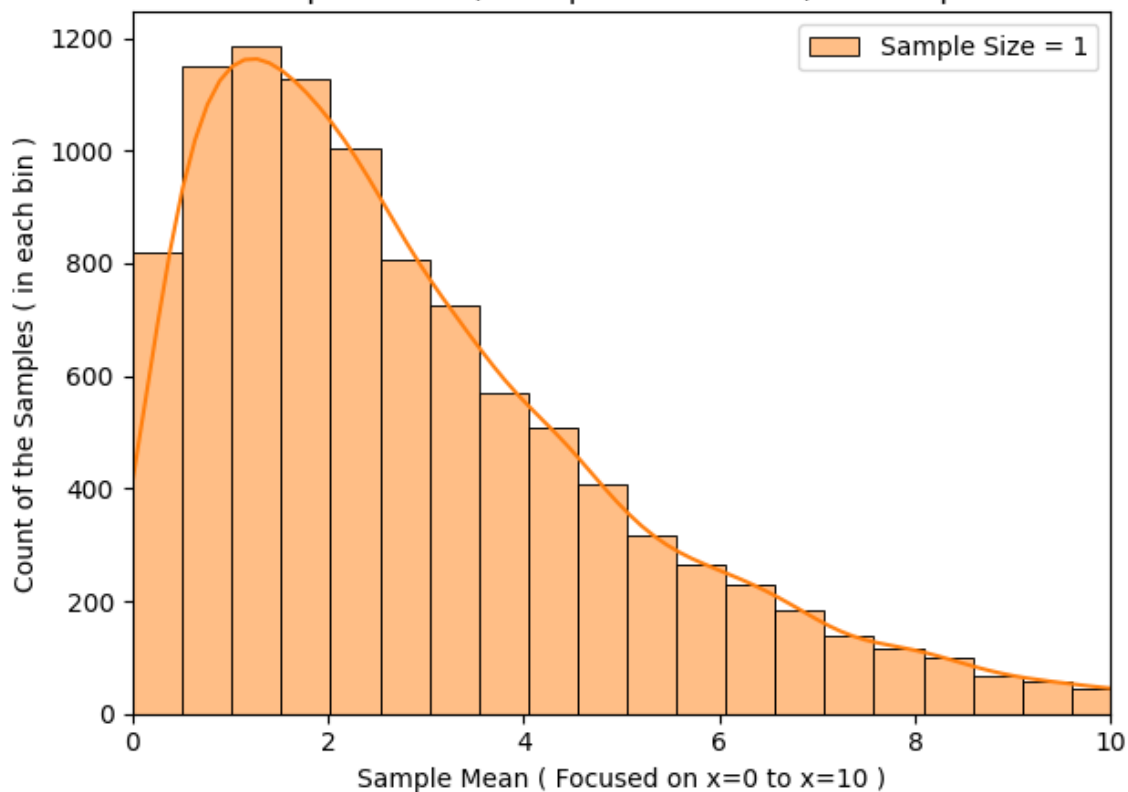
# Calculate sample means
sample_means = np.mean(samples, axis=0)

# Plot the Required Figures
sns.histplot(sample_means, kde=True, bins=50, label=f'Sample Size
= {sample_size}', color='C'+str(sample_size))
plt.legend()
plt.title(f'Distribution of Sample Means (Chi-square with df={df})
for samples drawn {sample_size} times')
plt.xlabel('Sample Mean ( Focused on x=0 to x=10 )')
plt.ylabel('Count of the Samples ( in each bin )')
plt.xlim(0, 10)
plt.tight_layout()
plt.show()

# Plot for samples drawn 1 times
first_plot(1)


```

Distribution of Sample Means (Chi-square with df=3) for samples drawn 1 times




```
# Plot for samples drawn 5 times  
first_plot(5)
```

Distribution of Sample Means (Chi-square with $df=3$) for samples drawn 5 times



```
# Plot for samples drawn 10 times  
first_plot(10)
```

Distribution of Sample Means (Chi-square with $df=3$) for samples drawn 10 times



2. The luminosity and redshift of galaxy clusters from XMM-BCS survey (details available at [arXiv:1512.01244](http://arxiv.org/abs/1512.01244)) can be downloaded <http://www.iith.ac.in/~shantanud/test.dat>. Plot the luminosity as a function of redshift on a log-log scale. By eye, do you think the datasets are correlated? Calculate the Spearman, Pearson and Kendall-tau correlation coefficients and the p-value for the null hypothesis.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import spearmanr, pearsonr, kendalltau
```

```

def second_plot()->None:

    # Load the data Stored on Local Machine in same folder
    data = pd.read_csv('galaxy_cluster_data.txt', sep='\s+',
header=None, names=['luminosity', 'redshift'])

    # Dropped the first row in the dataset as we have loaded the
    # dataset with different column names
    data = data.drop([0])

    # Convert columns to numeric (if needed)
    # I am getting TypeError: unsupported operand type(s) for +:
    'float' and 'str'
    # Hence I have used below function
    data['luminosity'] = pd.to_numeric(data['luminosity'],
errors='coerce')
    data['redshift'] = pd.to_numeric(data['redshift'],
errors='coerce')

    # Plotting the luminosity as a function of redshift on a log-log
    scale
    plt.scatter(x=data['redshift'], y=data['luminosity'], alpha=0.5,
c='C2', marker='x')
    plt.xscale('log')
    plt.yscale('log')
    plt.xlabel('Redshift (log scale)')
    plt.ylabel('Luminosity (log scale)')
    plt.title(' Redshift Vs Luminosity ')
    plt.show()


    # Calculate correlation coefficients and p-values
    spearman_corr, spearman_p_value = spearmanr(data['redshift'],
data['luminosity'])
    pearson_corr, pearson_p_value = pearsonr(data['redshift'],
data['luminosity'])
    kendall_corr, kendall_p_value = kendalltau(data['redshift'],
data['luminosity'])

    # Print the results
    print(f"\n\nSpearman correlation coefficient: {spearman_corr:.4f},
Spearman p-value: {spearman_p_value:.4f}")
    print(f"\n\nPearson correlation coefficient: {pearson_corr:.4f},
Pearson p-value: {pearson_p_value:.4f}")
    print(f"\n\nKendall Tau correlation coefficient:
{kendall_corr:.4f},    Kendall Tau p-value: {kendall_p_value:.4f}")

# Call the function to display plot
second_plot()

```

Before calculating the correlation coefficient, We observe
that there is a Linear-Positive Relationship among Redshift and
Luminosity



Spearman correlation coefficient: 0.6596, Spearman p-value: 0.0000

Pearson correlation coefficient: 0.5144, Pearson p-value: 0.0003

Kendall Tau correlation coefficient: 0.5030, Kendall Tau p-value:
0.0000

3. Wind speed data from the Swiss Wind Power data website can be found at <http://wind-data.ch/tools/weibull.php>. Using the data provided on the website, plot the probability distribution and overlay the best-fit Weibull distribution (with the parameters shown on the website). (20 points) (Hint : A on the website is same as λ , which was used in class to parameterize the Weibull distribution.)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import weibull_min

def third_plot()->None:
    # Load the data Stored on Local Machine in same folder
    data = pd.read_csv('wind_data.txt', sep='\t', header=None,
names=['Class', 'Frequency in %'])

    # Dropped the first row in the dataset as we have loaded the
# dataset with different column names
    data = data.drop([0])

    # Convert columns to numeric (if needed)
    # TypeError: ufunc 'isfinite' not supported for the input types,
and the inputs could
    # not be safely coerced to any supported types according to the
casting rule 'safe'
    # Hence I have used below function
    data['Class'] = pd.to_numeric(data['Class'], errors='coerce')
    data['Frequency in %'] = pd.to_numeric(data['Frequency in %'],
errors='coerce')

    # Weibull distribution parameters (replace these with the actual
parameters)
    shape = 2.0 # shape parameter (c or k)
    scale = 6.0 # scale parameter ( $\lambda$  or c)
    location = 0.0 # location parameter ( $\delta$ )
```



```

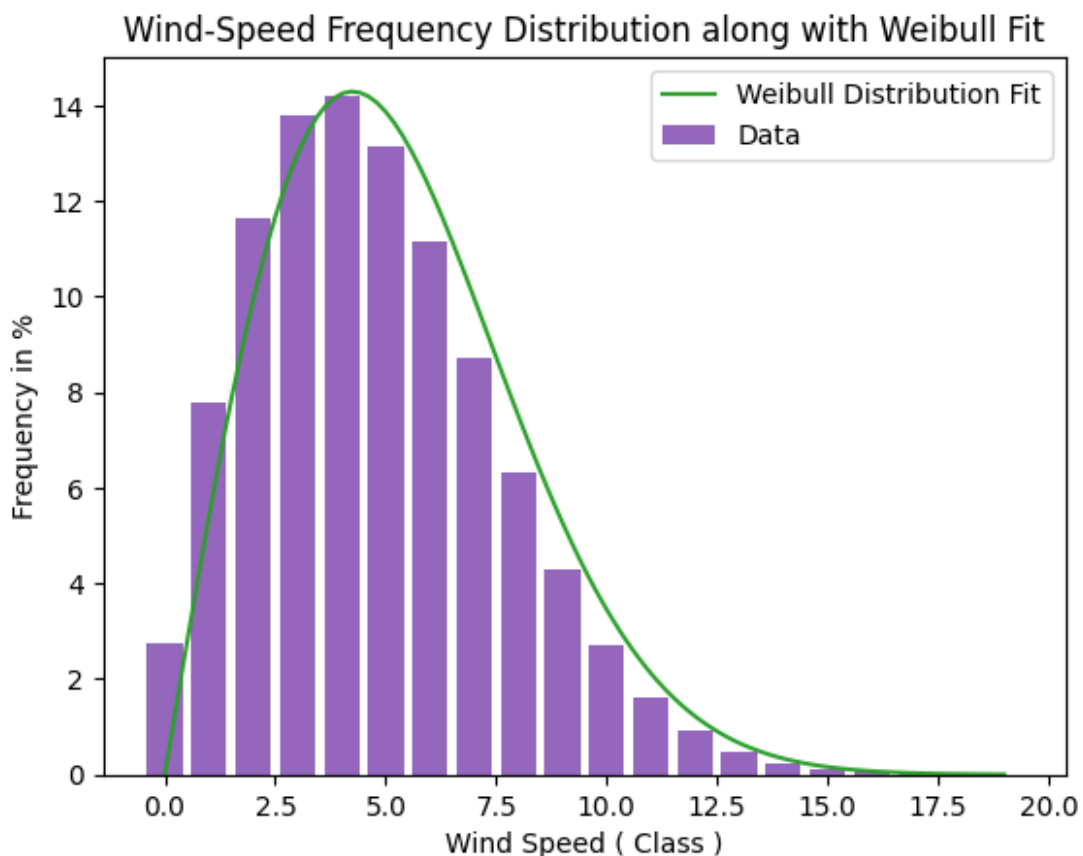
# Plot the data in the histogram for the % frequencies
plt.bar(data['Class'], data['Frequency in %'], label='Data',
color='C4')

# Overlay the Weibull distribution
x = np.linspace(0, max(data['Class']), 100)
y_fit = weibull_min.pdf(x, shape, loc=location, scale=scale)
plt.plot(x, y_fit * 100, 'C12', label='Weibull Distribution Fit')

plt.xlabel('Wind Speed ( Class )')
plt.ylabel('Frequency in %')
plt.title('Wind-Speed Frequency Distribution along with Weibull
Fit')
plt.legend()
plt.show()

# Call the function to display plot
third_plot()

```



4. Generate two arrays of size 1000 drawn from a Gaussian distribution of mean of zero and standard deviation of one. Calculate Pearson correlation coefficient and its p-value using scipy module. Also check if the p-value agrees with that calculated using the Student-t distribution.

```
import numpy as np
from scipy.stats import pearsonr, t

def fourth_plot(t_dis_tail: int=2)->None:

    np.random.seed(42) # Set the seed to draw the random samples

    # Generate two Numpy arrays of size 1000 drawn from a Gaussian
    # distribution
    gaussian_generated_one = np.random.normal(0, 1, 1000)
    gaussian_generated_two = np.random.normal(0, 1, 1000)

    # Calculate Pearson correlation coefficient and its p-value
    correlation_pearson, pearson_p_value =
    pearsonr(gaussian_generated_one, gaussian_generated_two)

    # Degrees of freedom for the Student-t distribution
    df = len(gaussian_generated_one) - 2 # for Pearson correlation

    # Calculate p-value using the Student-t distribution
    t_distribution_p_value = t_dis_tail * (1 -
    t.cdf(np.abs(correlation_pearson) * np.sqrt(df) / np.sqrt(1 -
    correlation_pearson**2), df))

    print(f"\n\nPearson correlation coefficient:
    {correlation_pearson:.4f}")
    print(f"\n\nP-value using scipy.stats.pearsonr:
    {pearson_p_value:.4f}")
    print(f"\n\nP-value using Student-t distribution:
    {t_distribution_p_value:.4f}")

    # check if the p-value agrees with that calculated using the
    # Student-t distribution.
    if np.isclose(pearson_p_value, t_distribution_p_value, rtol=1e-4):
        print("\n\nP-values agree.")
```

```
    else:  
        print("\n\nP-values do not agree.")  
# Taking the t-distribution tail as 1 ( Half )  
fourth_plot(t_dis_tail=1)
```

Pearson correlation coefficient: -0.0404

P-value using scipy.stats.pearsonr: 0.2018

P-value using Student-t distribution: 0.1009

P-values do not agree.

```
# Taking the t-distribution tail as 1 ( Half )  
fourth_plot(t_dis_tail=2)
```

Pearson correlation coefficient: -0.0404

P-value using scipy.stats.pearsonr: 0.2018

P-value using Student-t distribution: 0.2018

P-values agree.